

CS2204

Program Design and Data
Structures for Scientific
Computing

Announcements

- Project #2 is posted
 - We will be creating another class that allows us to represent & manipulate DNA strands
 - We will use a list as our underlying storage container
 - Our goal is to have a faster/smaller representation by utilizing the mutability of lists (as compared to the immutability of strings)
 - You should use mutation operations on your list whenever possible
 - Avoid operations that create new lists



Matrix Programming

Introduction

By Richard T. Guy



Copyright © Software Carpentry 2010

This work is licensed under the Creative Commons Attribution License

See <http://software-carpentry.org/license.html> for more information.

Studying patients with Babbage's Syndrome

How effective are available treatments (A, B, & C)?

	A	B	C
John	2.5	3.5	3.0
Mary	3.0	1.5	3.0
Zura	2.5	2.0	5.5

How similar are patients' responses?

Can we use similarity to recommend treatments?

We can answer these questions with matrix algebra.

How to implement them in software?


Option 1: write loops

- Makes programs many times longer than the corresponding mathematics
- And it's hard code to debug...
- ...and tune

Option 2: use libraries written in low-level, high-performance languages like Fortran and C

- Someone else has written, debugged, and tuned all the loops
- But the interface is... awkward

SUBROUTINE CAXPY(N,CA,CX,INCX,CY,INCY)



Option 3: use a high-level language like MATLAB
Or a library like Python's NumPy

- Present a *data-parallel* programming model
 - Operate on entire arrays at once
 - No loops!
- Hide details of optimizations
 - Particularly differences between machines
- All provide basically the same features
 - Often wrappers around the same underlying libraries

NumPy (<http://numpy.scipy.org>)

Provides MATLAB-style arrays for Python
And many other things

A data parallel programming model

- Write $\mathbf{x}^*\mathbf{A}*\mathbf{x.T}$ to calculate \mathbf{xAx}^T
- The computer takes care of the loops

All encapsulated in special objects called *arrays*

Create an array from a list

```
>>> import numpy
>>> vals = [3, 5, 7]
>>> arr = numpy.array(vals)
>>> arr
array([3, 5, 7])
```

Alternatively...



```
>>> from numpy import *
>>> vals = [3, 5, 7]
>>> arr = array(vals)
>>> arr
array([3, 5, 7])
```

Arrays are *homogeneous*

- I.e., all values have the same type

Allows values to be packed together

- Saves memory
- Faster to process



So what does this do?

```
>>> arr = numpy.array([1, 2.3])
```

So what does this do?

```
>>> arr = numpy.array([1, 2.3])
```

```
>>> arr
```

```
array([1., 2.3])
```



A float, not an int

If we give NumPy initial values of different types...

...it finds the most general type---in this case, float---and uses that.

You can specify a type at creation time:

```
>>> array([1, 2, 3, 4], dtype=float32)
```

```
array([ 1., 2., 3., 4.])
```



You can also specify the data type later

```
>>> a = array([1, 2, 3, 4], dtype=float32)
```

```
>>> a.astype(int)
```

```
array([ 1, 2, 3, 4])
```

```
>>> #the above returns a new array
```

You can also specify the data type later

```
>>> a = array([1, 2, 3, 4], dtype=float32)
```

```
>>> a.astype(int)
```

```
array([ 1, 2, 3, 4])
```

```
>>> #the above returns a new array
```

```
>>> a.dtype = int
```

```
>>> a
```

```
array([1065353216, 1073741824, 1077936128,
       1082130432])
```

```
>>> #the above changed the type field, causing the bits
      to be interpreted differently
```

Basic data types are:

bool

int

int8

int16

int32

int64

uint[8,16,32,64]

float

float[16,32,64,128]

complex

complex[64,128]

Many other ways to create arrays

```
>>> z = numpy.zeros((2, 3))
```

```
>>> z
```

```
array([[0., 0., 0.],  
       [0., 0., 0.]])
```

- Type is **float** unless something else specified
- The 'zeros' function takes a tuple specifying array dimensions

Many other ways to create arrays

```
>>> z = numpy.zeros((2, 3))
```

```
>>> z
```

```
array([[0., 0., 0.],  
       [0., 0., 0.]])
```

- Type is **float** unless something else specified
- The 'zeros' function takes a tuple specifying array dimensions

What do these do?

```
>>> block = numpy.ones((4, 5))
```

```
>>> mystery = numpy.identity(4)
```


Can create arrays without filling in values

```
>>> x = numpy.empty((2, 2))
```

```
>>> x
```

```
array([[3.82265e-297, 4.94944e+173],  
       [1.93390e-309, 1.00000e+000]])
```

"Values" will be whatever bits were in memory

Should not be used without being initialized

When is this useful?

As with everything, assigning creates alias: does *not* copy data

```
>>> first = numpy.ones((2, 2))
```

```
>>> first
```

```
array([[1., 1.],  
       [1., 1.]])
```

```
>>> second = first
```

```
>>> second[0, 0] = 9
```

```
>>> first
```

```
array([[9., 1.],  
       [1., 1.]])
```

As with everything, assigning creates alias: does *not* copy data

```
>>> first = numpy.ones((2, 2))
```

```
>>> first
```

```
array([[1., 1.],
       [1., 1.]])
```

```
>>> second = first
```

```
>>> second[0, 0] = 9
```

```
>>> first
```

```
array([[9., 1.],
       [1., 1.]])
```

Not : **second[0][0]**



Use the **array.copy** method to make a copy

```
>>> first
```

```
array([[1., 1.],  
       [1., 1.]])
```

```
>>> second = first.copy()
```

```
>>> second[0, 0] = 9
```

```
>>> first
```

```
array([[1., 1.],  
       [1., 1.]])
```



Arrays also have properties

```
>>> first
```

```
array([[1., 1.],  
       [1., 1.]])
```

```
>>> first.shape
```

```
(2, 2)
```

```
>>> block = numpy.zeros((4, 7, 3))
```

```
>>> block.shape
```

```
(4, 7, 3)
```


Arrays also have properties

```
>>> first
```

```
array([[1., 1.],  
       [1., 1.]])
```

```
>>> first.shape
```

```
(2, 2)
```

← Not a method call

```
>>> block = numpy.zeros((4, 7, 3))
```

```
>>> block.shape
```

```
(4, 7, 3)
```

Arrays also have properties

```
>>> first
```

```
array([[1., 1.],  
       [1., 1.]])
```

```
>>> first.shape
```

```
(2, 2)
```

Not a method call



```
>>> block = numpy.zeros((4, 7, 3))
```

```
>>> block.shape
```

```
(4, 7, 3)
```

Consistent



array.size is the total number of elements

>>> first.size

4 ← 2×2

>>> block.size

84 ← $4 \times 7 \times 3$

Reverse on all axes with **array.transpose**

```
>>> first = numpy.array([[1, 2, 3],  
                        [4, 5, 6]])
```

```
>>> first.transpose()
```

```
array([[1, 4],  
       [2, 5],  
       [3, 6]])
```

```
>>> first  
array([[1, 2, 3],  
       [4, 5, 6]])
```

this creates an alias that appears to have the values stored differently; does not actually copy all the data

Flatten arrays using **array.ravel**

```
>>> first = numpy.zeros((2, 2, 2))
```

```
>>> second = first.ravel()
```

```
>>> second.shape
```

```
(8,)
```

this creates a one-dimensional alias for the original data

Think about the 2×4 array A:

```
>>> A
```

```
array([[1, 2, 3, 4],  
       [5, 6, 7, 8]])
```

‘A’ looks 2-dimensional

But computer memory is 1-dimensional
Must decide how to lay out values

Row-major order concatenates the rows

- Used by C and Python

1	2	3	4
5	6	7	8

Logical

1	2	3	4	5	6	7	8
---	---	---	---	---	---	---	---

Physical

Column-major order concatenates the columns

- Used by Fortran and MATLAB

1	2	3	4
5	6	7	8

Logical

1	5	2	6	3	7	4	8
---	---	---	---	---	---	---	---

Physical

No difference in usability or performance¹ ...
...but causes headaches when passing data from
one language to another
(Just like 0-based vs. 1-based indexing)

1: Performance can be affected if you write your own loops over large matrices;
you want to process the matrix in physical order if possible

Can *reshape* arrays in many other ways

```
>>> first = numpy.array([1, 2, 3, 4, 5, 6])
```

```
>>> first.shape
```

(6,) ← Tuple with 1 element

```
>>> second = first.reshape(2, 3)
```

```
>>> second
array([[1, 2, 3],
       [4, 5, 6]])
```

↑
Not packed into a tuple

Also aliases the data

New shape must have same size as old

```
>>> first = numpy.zeros((2, 2))
```

```
>>> first.reshape(3, 3)
```

**ValueError: total size of new array must
be unchanged**

Cannot possibly work because it is just creating an
alias for the existing data

Change physical size using **array.resize**

```
>>> block
```

```
array([[ 10, 20, 30],
       [110, 120, 130],
       [210, 220, 230]])
```

```
>>> block.resize(2, 2) #may or may not work, depends on
                        python version
```

```
>>> block
```

```
array([[ 10, 20],
       [110, 120]])
```

This did not work for me; see next slide...

Change physical size using **resize**

```
>>> block
```

```
array([[ 10,  20,  30],  
       [110, 120, 130],  
       [210, 220, 230]])
```

```
>>>resize(block,(2, 2))
```

```
array([[ 10,  20],  
       [30, 110]])
```

Review:

- Arrays are blocks of homogeneous data
- Most operations create aliases
- Can be reshaped (size remains the same)
- Or resized